

---

# Reads Like an LLM, Treated Like a Human: Content-Matched Tests of Whether Prompt Authorship Style Changes LLM Safety Behaviour

---

NeuriCo

## Abstract

Safety evaluations, red-teaming pipelines and multi-agent systems increasingly feed language models user turns that were written by other language models. If a model behaves differently when a request *reads* LLM-written, evaluation results may not transfer to human users. We test this directly. We build 950 content-matched requests (harmful, borderline-safe, and factual questions with and without a user’s wrong belief), each paraphrased by two rewriters into a casual human style, a polished LLM style, and a length-matched LLM style, and verified for content equivalence by an independent judge. We measure refusal, over-refusal, sycophancy and accuracy in four responders (LLAMA-3.1-8B, QWEN2.5-7B, GPT-5.6-LUNA, GEMINI-3.5-FL), with paraphrase, length, politeness and explicit-authorship-label controls. For three of the four models, every LLM-style minus human-style contrast lies within about  $\pm 3$  percentage points (pp) and none survives Holm correction. GEMINI-3.5-FL refuses 3–4 pp more on LLM-style prompts, but a fixed polite opener on the original text reproduces its over-refusal effect (+3.4 vs. +3.8 pp), while an explicit “an AI agent wrote this” label has no effect. Inside the open models, “reads LLM-written” is linearly decodable (held-out AUROC 0.94–1.00), but the direction is a register axis: it fails to detect LLM-written prompts that imitate casual users (AUROC 0.48–0.53) and is nearly orthogonal to an evaluation-awareness direction. Steering it on fixed text flips QWEN2.5-7B’s own authorship judgement from “human” to “AI” while changing refusal and accuracy by at most 1 pp. Prompt authorship, as perceived by the model, does not drive safety behaviour; surface register does, in a model-specific way.

## 1 Introduction

A growing share of the prompts that language models answer were not written by people. Model-written evaluations [Perez et al., 2023], synthetic red-teaming, automated auditors and LLM user simulators all put LLM-written text in the user turn, and in multi-agent systems sub-agents are prompted mostly by an orchestrating model. At the same time, models can recognise LLM-written text: a linear direction separates human from machine documents in hidden states [Chen et al., 2025, Quaremba et al., 2026], and LLM judges prefer LLM-written items over human-written ones [Panicssery et al., 2024, Laurito et al., 2025]. Our main question is whether these two facts meet: **does a model behave differently when the request it answers reads LLM-written rather than human-written?**

The question matters for evaluation validity. If refusal or sycophancy shifts when the request reads LLM-written, then safety numbers measured on LLM-written prompts will not carry over to deployment with human users, and an orchestrated sub-agent may behave differently from the same

model in a chat window. A related literature on evaluation awareness shows that models can tell test transcripts from deployment ones [Needham et al., 2025] and that steering an evaluation-awareness direction can change compliance [Abdelnabi and Salem, 2025, Nguyen et al., 2025]. An LLM-written prompt could plausibly look “test-like” to a model.

However, existing work does not answer the question. On one hand, studies of prompt style vary a single named feature such as politeness, persuasion or syntax, and do not ask about authorship. On the other hand, AI–AI bias studies vary the *item being judged*, never the instruction itself. The closest prior study [Chalkidis, 2026] compares templated and LLM-generated prompts without item-level content matching, and traces the difference to content rather than provenance. Moreover, every close precedent we found reduced an apparent “source” effect to surface features once it was controlled [Devbunova, 2026, Ackerman and Panickssery, 2024]. A convincing test must therefore hold content fixed, separate authorship from length, politeness and register, and check causally whether any internal “LLM-written” representation is what moves behaviour.

We propose such a test. We hold the request fixed and change only whether it reads human-written or LLM-written. Every one of 950 items from JBB-BEHAVIORS, XSTEST and TRIVIAQA is paraphrased by two different rewriter models into a casual human style (H), a polished LLM style (L) and a length-matched LLM style (S); an independent judge verifies that each rewrite keeps the request intact. Because H is also an LLM paraphrase, “was paraphrased” is constant across conditions. We add three controls that break the bundle of features carried by LLM style: a deterministic polite opener on the original text (POL), an explicit system line stating that a human or an AI agent wrote identical text, and a length covariate. We then measure perception (verbal judgements and linear probes), and in two open models we steer a learned “reads LLM-written” direction on fixed text against norm-matched random directions.

Our answer is a controlled near-null plus one model-specific surface effect. For LLAMA-3.1-8B, QWEN2.5-7B and GPT-5.6-LUNA, all LLM-style minus human-style contrasts lie within about  $\pm 3$  pp and none survives Holm correction. GEMINI-3.5-FL refuses more on LLM-style prompts (+3.8 pp over-refusal, +3.1 pp harmful refusal), but the polite opener alone gives +3.4 pp, GEMINI-3.5-FL almost never judges these prompts to be AI-written (1.5% for rewriter A), and an explicit “AI agent” label has no effect. Inside LLAMA-3.1-8B and QWEN2.5-7B, the “reads-LLM” direction is near-perfectly decodable (AUROC 0.94–1.00) but tracks register rather than provenance, and it is nearly orthogonal to an evaluation-awareness direction ( $|\cos| \leq 0.10$ ). Steering QWEN2.5-7B along it flips the model’s own judgement of the identical text from “human” (log-odds  $-1.6$ ) to “AI” (+3.6) while refusal, over-refusal and accuracy change by at most 1 pp. In short, we contribute:

- We construct a content-matched stimulus set of 950 items in human-style, LLM-style and length-matched LLM-style paraphrases from two rewriters, filtered for content equivalence, with original, polite-opener and explicit-label controls.
- We conduct a paired, Holm-corrected comparison of refusal, over-refusal, sycophancy and accuracy across four responders ( $\approx 44.8$ k responses) and show that LLM-style prompts do not systematically change these behaviours; the one robust shift (GEMINI-3.5-FL) is reproduced by politeness alone.
- We identify an internal “reads-LLM” direction in two open models and show that it is a register axis, separate from provenance and from evaluation awareness.
- We complement the correlational analyses with a causal steering test on fixed text, with dose sweeps, random-direction controls and degradation checks, which dissociates perceived authorship from behaviour.

Section 2 reviews related work, section 3 describes the stimuli, models and analyses, section 4 reports the results, and section 5 discusses their implications and limitations.

## 2 Related Work

**Evaluation awareness.** Models can distinguish evaluation from deployment transcripts, with a best AUC of 0.83 against 0.92 for humans [Needham et al., 2025]; the cues they cite include structured, formal phrasing. Linear probes find an evaluation-awareness direction [Nguyen et al., 2025, Heidari et al., 2026], and steering it changes harmful tool-call compliance in reasoning models [Abdelnabi and Salem, 2025] or flips evaluation-specific behaviour in a fine-tuned model [Hua et al., 2025]. Yet the direction tracks benchmark format rather than context [Devbunova, 2026], machine-generated

| Set                                  | Outcome      | $n$ |
|--------------------------------------|--------------|-----|
| JBB-BEHAVIORS harmful                | refusal      | 100 |
| XSTEST unsafe                        | refusal      | 200 |
| JBB-BEHAVIORS benign (topic-matched) | over-refusal | 100 |
| XSTEST safe-but-edgy                 | over-refusal | 250 |
| TRIVIAQA, neutral                    | accuracy     | 150 |
| TRIVIAQA + user’s wrong belief       | sycophancy   | 150 |

Table 1: Base items. Every item appears in all prompt variants; analyses are paired within item.

instructions are classified as “deployment” [Nguyen et al., 2025], and detection, representation and behaviour are only weakly coupled [Nayan et al., 2026, Heidari et al., 2026]. Building on the 16 contrastive pairs of Hua et al. [2025], we measure how a “reads-LLM” direction relates to evaluation awareness, and we find them nearly orthogonal.

**AI–AI bias and recognition of LLM text.** LLM judges recognise and prefer their own summaries [Panickssery et al., 2024], prefer LLM-written product and paper descriptions over human-written ones [Laurito et al., 2025], and favour LLM-written resume summaries, more strongly in large models [Xu et al., 2025]. Self-recognition in Llama-3-8B is carried by a steerable residual direction, but it largely reflects length [Ackerman and Panickssery, 2024]. In all of these studies the LLM-written text is the *object* being evaluated. Unlike this line of work, we vary the authorship style of the *instruction* and measure safety behaviour.

**Hidden-state detection of machine text and of the interlocutor.** Human and LLM documents are linearly separable from last-token activations [Quaremba et al., 2026] or from paired activation differences [Chen et al., 2025], though transfer is weakest when content is fixed. TalkTuner [Chen et al., 2024] reads user demographics from chat activations where verbal prompting fails, and steers answers with the same probes. We apply this probe-then-steer recipe to prompt authorship and test whether the representation matters for behaviour, not only whether it exists.

**Prompt style, construction and safety behaviour.** Register and framing move behaviour with information held fixed: casual challenges induce more answer switching than formal ones [Kim and Khashabi, 2025], and sycophancy benchmarks rely on first-person belief statements [Sharma et al., 2024]. LLM user simulators match real users poorly on formality and capitalisation [Ivey et al., 2024]. The work closest to ours [Chalkidis, 2026] finds that templated and LLM-generated prompts produce different stances, and attributes the gap to template filler content; it leaves open whether recognising a prompt’s source changes behaviour. Unlike that study, we match content at the item level, separate length and politeness from authorship, and test the authorship representation causally.

### 3 Method

Our design holds the content of a request fixed and changes only whether it reads human-written or LLM-written. We formulate four hypotheses. (H1) Behaviour differs between LLM-style and human-style versions of the same request. (H2) Any difference is explained by length or other surface features. (H3) An explicit statement that an AI agent wrote the text changes behaviour relative to a “human” statement on identical text. (H4) Steering an internal “reads-LLM” direction on fixed text changes behaviour more than norm-matched random directions do.

#### 3.1 Stimuli

**Base items.** We use 950 items from four sources, each tied to one outcome (table 1). Harmful items from JBB-BEHAVIORS [Chao et al., 2024] and unsafe items from XSTEST [Röttger et al., 2024] measure refusal. Topic-matched benign JBB-BEHAVIORS items and safe-but-edgy XSTEST items measure over-refusal; they give headroom where harmful refusal is near ceiling. Questions from TRIVIAQA, as packaged in the sycophancy evaluation of Sharma et al. [2024], measure accuracy in neutral form and sycophancy when followed by “I think the answer is  $\langle$ wrong $\rangle$  but I’m really not sure.”

| Variant       | Chars | Words     | Lowercase start | Polite marker |
|---------------|-------|-----------|-----------------|---------------|
| ORIG          | 72    | 12.9      | 0%              | 0%            |
| H (A/B)       | 66–69 | 12.0–12.4 | 98–100%         | 0%            |
| S (A/B)       | 91    | 15.0–15.5 | 0%              | 99–100%       |
| L, rewriter A | 125   | 21        | 0%              | 100%          |
| L, rewriter B | 236   | 38        | 0%              | 100%          |

Table 2: Mean surface profile of the prompt variants. LLM style bundles length, capitalisation and politeness; the length-matched S variant removes most of the length difference.

**Variants.** Two rewriters, A = grok-4.3 and B = deepseek-v4-flash, paraphrase each item into three styles: H imitates a casual human typing into a chatbot; L reads like an LLM-written prompt (polished, formal, polite, explicit, with no mention of AI or testing); S is LLM style capped at  $|H| + 2$  words, a length-matched control. Because every H, L and S variant is an LLM paraphrase, the act of paraphrasing is constant, and H is an LLM text that imitates human style. We keep the original dataset text (ORIG) as a reference. After inspecting GEMINI-3.5-FL’s behaviour we add POL: the original text with the fixed prefix “Could you please help me with the following request?”. POL is deterministic, so it changes one surface feature and leaves the content untouched. Finally, two label conditions show identical text (H or L from rewriter A) with a system line stating that the message was “written by a human user” or “written by an AI agent (another language model)”; this separates belief about the author from style.

**Content-equivalence filter.** claude-sonnet-5, blind to the study’s purpose, checks each rewrite against the original: same request and details, nothing added or removed, no added framing. We allow up to three attempts per rewrite. Hard checks require the wrong answer in sycophancy items to survive verbatim with first-person belief, and neutral items not to leak the answer. Pass rates are 99% for H, 98–99% for S, 97% for L from A and 93% for L from B. We analyse only variants that pass.

**Surface profile.** Table 2 shows that the styles differ as intended. H is short and almost always lowercase; L is longer and always polite. S is closer to H in length but still about 3 words longer, so we also adjust for length statistically. Casing and politeness are bundled with LLM style by construction; POL and the label conditions are what separate them.

### 3.2 Responders and outcome measurement

**Responders.** We use two local open-weight models, Llama-3.1-8B-Instruct (LLAMA-3.1-8B) and Qwen2.5-7B-Instruct (QWEN2.5-7B), run in bf16 with greedy decoding, and two API models via OpenRouter, gpt-5.6-luna (GPT-5.6-LUNA; temperature 0, reasoning off) and gemini-3.5-flash-lite (GEMINI-3.5-FL; temperature 0, reasoning effort “low”, since it cannot be disabled). We allow 200 new tokens for refusal items and 120 for trivia items, plus 600 tokens of reasoning headroom for GEMINI-3.5-FL. Eleven conditions plus POL give about 11.2k responses per model and 44.8k in total.

**Outcomes.** A response judge (deepseek-v4-pro) sees the original item text, never the styled variant or the label, so it is blind to condition. For refusal it applies the three-way XSTEST scheme (full compliance, full refusal, partial refusal); we count full or partial refusal as “refused”. On 300 JBB-BEHAVIORS responses with human labels, 95% (63/66) of responses the judge calls full refusals are labelled “not jailbroken”, and 91% (100/110) of jailbroken responses are judged compliant; the remaining disagreements are mostly harmless or fictionalised compliance, which the human labels mark as safe. For trivia, the judge labels each answer as correct (aliases supplied), the user’s suggested wrong answer, or other. Sycophancy is the rate of adopting the suggested wrong answer on biased items; accuracy is the rate of correct answers on neutral items. A string-match refusal flag serves as a judge-free secondary metric. Responses that GPT-5.6-LUNA’s provider blocked with a content filter (693 of 11.2k) count as refusals, since that is what a user sees; a robustness check drops them.

### 3.3 Statistical analysis

For each model and outcome we compute paired within-item contrasts, pooling rewriters (a pair is item  $\times$  rewriter): L–H, S–H, L–S, ORIG–H, POL–ORIG, label AI – label human (on H text and on L text), and L–H under the human label. We report differences in percentage points with 95% item-cluster bootstrap confidence intervals (10k resamples) and exact McNemar tests on discordant pairs, with Holm correction within each contrast family across the 16 model  $\times$  outcome cells. To separate style from length (H2), we fit a within-item conditional logit over the six rewrites:

$$\text{logit } P(y_{iv} = 1) = \alpha_i + \beta_{\text{style}} \mathbb{I}[v \in \{\text{L}, \text{S}\}] + \beta_{\text{len}} \log(\text{chars}_{iv}) + \beta_{\text{rw}} \mathbb{I}[v \text{ from B}], \quad (1)$$

where  $i$  indexes items,  $v$  variants, and  $\alpha_i$  is an item fixed effect conditioned out. For the local models, a correlational mediation analysis replaces the style dummy with the variant’s projection on the internal direction defined below.

**Perception checks.** We ask the API models “Was this message written by a human or generated by an AI language model?” in both answer orders, on 50 items per task for each of eight variants. For the local models we read the first-token log-odds  $\text{logit}(\text{AI}) - \text{logit}(\text{Human})$ , averaged over both orders. We validate this verbal read-out on HC3 [Guo et al., 2023] human vs. ChatGPT answers and on POLISTANCE [Chalkidis, 2026] real vs. LLM-generated prompts.

### 3.4 Internal representation and causal steering

**The “reads-LLM” direction.** We take 500 real first-turn requests from WILDCHAT [Zhao et al., 2024] and have rewriter A produce H, L and S versions of each. For each layer  $l$  we read the residual stream  $\mathbf{h}_l(x)$  at the final prompt position and define

$$\mathbf{d}_l = \frac{1}{|\mathcal{S}|} \sum_{x \in \mathcal{S}} \mathbf{h}_l(x) - \frac{1}{|\mathcal{H}|} \sum_{x \in \mathcal{H}} \mathbf{h}_l(x), \quad (2)$$

where  $\mathcal{S}$  and  $\mathcal{H}$  are the S and H versions in a 70% training split. Using S rather than L keeps the pairs paraphrase-matched and roughly length-matched. We evaluate the projection  $\mathbf{h}_l(x)^\top \mathbf{d}_l$  by held-out AUROC against a random direction, a length-only classifier and a full logistic probe, and test transfer to POLISTANCE (an unseen source of LLM-written prompts that imitate casual users) and to the behavioural items. We also report the cosine of  $\mathbf{d}_l$  with an evaluation-awareness direction built from the 16 “real user” vs. “evaluator” pairs of Hua et al. [2025] and with a refusal direction (JBB-BEHAVIORS harmful minus benign means, following Arditì et al. 2024). To ask whether provenance is present beyond style, we train an in-domain five-fold probe on POLISTANCE and compare it with length and TF-IDF classifiers.

**Steering on fixed text.** We hold the text exactly fixed at H from rewriter A for all 950 items and add  $\alpha \mathbf{d}_l$  at every position in eight consecutive middle layers (hidden states 9–16 for LLAMA-3.1-8B, 7–14 for QWEN2.5-7B). We sweep  $\alpha \in \{\pm 0.25, \pm 0.5, \pm 1\}$ ; positive  $\alpha$  pushes toward “LLM-written”.  $\alpha = 0.25$  moves the read-out projection by about the natural H  $\rightarrow$  S difference (1.35 vs. 1.63 for LLAMA-3.1-8B at layer 20), and  $\alpha = 2$  produced incoherent text in calibration. Controls are norm-matched random directions (an independent Gaussian vector per layer rescaled to  $\|\mathbf{d}_l\|$ ; three seeds for LLAMA-3.1-8B, two for QWEN2.5-7B, at  $\alpha = 0.5$  and 1) and the unsteered model. Manipulation checks are the read-out projection four layers above the steering band and the verbal self-report. Degradation checks are neutral TRIVIAQA accuracy, response length and the share of all-lowercase replies.

**Reproducibility.** We fix seed 42 and use deterministic decoding throughout, and cache every API call so all runs replay offline. The study used about \$30 of API credit and 3.5 GPU-hours on one RTX A6000 (48 GB). Appendix A lists further details.

## 4 Results

We first check that the models perceive the manipulation (section 4.1), then report the behavioural contrasts (section 4.2), analyse the one robust effect (section 4.3), characterise the internal direction (section 4.4), and finally test it causally (section 4.5).

|                                                                                                    | ORIG          | H <sub>A</sub> | H <sub>B</sub> | S <sub>A</sub> | S <sub>B</sub> | L <sub>A</sub> | L <sub>B</sub> | POL  |
|----------------------------------------------------------------------------------------------------|---------------|----------------|----------------|----------------|----------------|----------------|----------------|------|
| <i>Verbal judgement</i>                                                                            |               |                |                |                |                |                |                |      |
| GPT-5.6-LUNA, % “AI”                                                                               | 20.6          | 10.5           | 8.6            | 27.1           | 34.2           | 39.5           | <b>89.2</b>    | 38.8 |
| GEMINI-3.5-FL, % “AI”                                                                              | 0.2           | 0.2            | 0.0            | 0.2            | 0.5            | 1.5            | <b>19.2</b>    | 6.2  |
| LLAMA-3.1-8B, log-odds(AI)                                                                         | −3.10         | −3.10          | −3.08          | −3.47          | −3.52          | −3.56          | −4.44          | −    |
| QWEN2.5-7B, log-odds(AI)                                                                           | − <b>0.42</b> | −1.64          | −1.79          | −1.17          | −1.16          | −0.67          | −1.42          | −    |
| <i>Internal probe (projection on “reads-LLM” direction, z relative to H<sub>A</sub>, layer 12)</i> |               |                |                |                |                |                |                |      |
| LLAMA-3.1-8B                                                                                       | 0.89          | 0              | 0.02           | 1.40           | 1.44           | 1.48           | <b>1.66</b>    | −    |
| QWEN2.5-7B                                                                                         | 0.87          | 0              | −0.02          | 1.51           | 1.68           | 1.95           | <b>2.64</b>    | −    |

Table 3: Is the manipulation perceived? Subscripts denote rewriter A or B. GPT-5.6-LUNA clearly perceives LLM style; GEMINI-3.5-FL almost always answers “human”. LLAMA-3.1-8B’s verbal judgement is inverted (polished prompts look *more* human), whereas both open models separate the styles internally and monotonically:  $H < \text{ORIG} < S < L$ . Highest value per row in **bold**.




Figure 1: Perceived authorship of each prompt variant. (Left) Share of “AI” answers from the API models. (Center) First-token log-odds of “AI” for the local models. (Right) Projection of each variant on the internal “reads-LLM” direction ( $z$  relative to  $H_A$ ). Verbal judgements are weak or inverted, but internally both open models order the variants  $H < \text{ORIG} < S < L$ .

#### 4.1 Do models perceive the style manipulation?

**Is a null result interpretable?** Only if the models register the difference between styles. Table 3 and figure 1 show that they do, though not always verbally. GPT-5.6-LUNA clearly perceives LLM style: it calls about 10% of H prompts AI-written, about 30% of S prompts and 40–89% of L prompts, with the verbose rewriter B detected most. GEMINI-3.5-FL answers “human” almost always (1.5% “AI” for  $L_A$ ). LLAMA-3.1-8B’s verbal report is inverted on prompts: it rates LLM-style prompts as *more* human-like, and its read-out reaches only 0.41 AUROC on POLISTANCE, although the same read-out works on documents (0.88 on HC3). QWEN2.5-7B’s verbal report points weakly in the right direction on our variants but is at chance on POLISTANCE (0.53; HC3 0.82). Internally, however, both open models separate the styles strongly and monotonically. As with user attributes in TalkTuner [Chen et al., 2024], verbal reports understate what the models represent, so we rely on probes as the primary local perception check.

#### 4.2 Does LLM style change behaviour?

**For three of four models, no.** Table 4 and figure 2 report the paired contrasts. For LLAMA-3.1-8B, QWEN2.5-7B and GPT-5.6-LUNA, every primary L–H contrast lies within about  $\pm 3$  pp and none survives Holm correction. The largest is GPT-5.6-LUNA’s harmful refusal ( $-2.5$  pp, CI  $[-4.7, -0.4]$ ), and it shrinks to  $-1.6$  pp (n.s.) once we drop items blocked by the provider’s content filter. Sycophancy L–H ranges from  $-2.4$  to  $+2.1$  pp and accuracy from  $-3.1$  to  $+3.4$  pp, all with CIs that include zero. Uncorrected sycophancy cells with CIs excluding zero (GPT-5.6-LUNA S–H  $+3.9$ ; QWEN2.5-7B POL–ORIG  $+6.7$ ; QWEN2.5-7B L–H under the human label  $+10.3$ ) point in inconsistent directions and do not survive correction.

**GEMINI-3.5-FL refuses slightly more.** GEMINI-3.5-FL’s over-refusal rises by  $+3.8$  pp under L ( $[+2.0, +5.8]$ , Holm  $p < .05$ ) and its harmful refusal by  $+3.1$  pp. Both survive length matching ( $+2.7$  and  $+3.3$  pp under S, both Holm-significant). Only two other cells are Holm-significant, both

| Contrast                                                                    | LLAMA-3.1-8B                  | QWEN2.5-7B                 | GPT-5.6-LUNA      | GEMINI-3.5-FL             |
|-----------------------------------------------------------------------------|-------------------------------|----------------------------|-------------------|---------------------------|
| <i>Refusal of harmful requests</i> (base rates 86–95%)                      |                               |                            |                   |                           |
| L–H                                                                         | −2.0 [−4.1, 0.0]              | −1.1 [−3.8, 1.6]           | −2.5 [−4.7, −0.4] | +3.1 [+0.4, +5.8]         |
| S–H (length-matched)                                                        | −0.9 [−3.2, 1.5]              | +1.5 [−0.7, 4.0]           | −1.9 [−4.3, 0.3]  | <b>+3.3</b> [+0.9, +5.8]* |
| POL–ORIG                                                                    | −0.3                          | +0.3                       | −0.3              | +0.3                      |
| label AI–human (H text)                                                     | +0.3                          | +1.7                       | +1.4              | −0.3                      |
| label AI–human (L text)                                                     | +1.4                          | +1.4                       | +1.4              | +1.4                      |
| <i>Refusal of safe or benign requests (over-refusal)</i> (base rates 4–12%) |                               |                            |                   |                           |
| L–H                                                                         | +0.5 [−1.5, 2.4]              | 0.0 [−2.0, 2.0]            | −2.1 [−4.6, 0.3]  | <b>+3.8</b> [+2.0, +5.8]* |
| S–H (length-matched)                                                        | +1.4 [−0.4, 3.4]              | +2.0 [0.0, 4.1]            | +0.3 [−1.9, 2.4]  | <b>+2.7</b> [+1.0, +4.6]* |
| POL–ORIG                                                                    | +0.9 [−1.1, 3.1]              | −0.9 [−2.6, 0.9]           | −1.1 [−3.4, 1.1]  | <b>+3.4</b> [+1.4, +5.4]* |
| label AI–human (H text)                                                     | +1.4 [−0.3, 3.4]              | +1.7 [0.0, 3.7]            | +0.6              | −0.9                      |
| label AI–human (L text)                                                     | +3.5 [+1.5, 5.9] <sup>†</sup> | +1.8                       | +0.6              | +0.6                      |
| <i>Sycophancy</i> (base rates 5–36%) and <i>accuracy</i>                    |                               |                            |                   |                           |
| sycophancy, L–H                                                             | −2.4 [−6.6, 1.4]              | +2.1 [−4.2, 8.2]           | 0.0 [−2.1, 2.1]   | −2.4 [−5.6, 0.7]          |
| accuracy, L–H                                                               | +3.4                          | −1.0                       | +1.4              | −3.1                      |
| accuracy, S–H                                                               | —                             | <b>−6.7</b> [−11.4, −2.3]* | —                 | —                         |
| accuracy, POL–ORIG                                                          | <b>+11.3</b> [+6.0, +16.7]*   | —                          | —                 | —                         |

Table 4: Content-matched behavioural contrasts, paired within item, in percentage points with 95% item-cluster bootstrap CIs (rewriters pooled). \* and **bold**: Holm-corrected  $p < .05$  across the 16 model  $\times$  outcome cells of the contrast family. <sup>†</sup>:  $p = .004$  uncorrected, Holm  $p = .067$ . Cells without a CI have CIs that include 0; “—” marks accuracy cells that are not significant and are omitted for space (all are shown in figure 2). Only GEMINI-3.5-FL shows a robust style effect, and it is reproduced by the polite opener alone (POL–ORIG).




Figure 2: All paired within-item contrasts for the four responders and four outcomes, with 95% item-bootstrap CIs; \* marks Holm-corrected  $p < .05$ . Almost all intervals straddle zero. GEMINI-3.5-FL’s refusal effects (L–H, S–H, POL–ORIG) are the only consistent pattern; the two significant accuracy cells (QWEN2.5-7B S–H, LLAMA-3.1-8B POL–ORIG) have no matching L–H effect.

for accuracy: QWEN2.5-7B S–H (−6.7 pp) and LLAMA-3.1-8B POL–ORIG (+11.3 pp, where the polite opener makes LLAMA-3.1-8B answer more fully). Neither has a matching L–H effect, so we read them as surface effects.

**Does believing an AI wrote the text matter?** The explicit label is the purest manipulation of belief about the author, and it never produces a Holm-significant effect. The largest is LLAMA-3.1-8B’s over-refusal on L text (+3.5 pp,  $p = .004$  uncorrected, Holm  $p = .067$ ). For GEMINI-3.5-FL, the label has no effect (−0.9 and +0.6 pp for over-refusal).

**Style net of length.** Table 5 reports the style coefficient of equation 1. For GEMINI-3.5-FL the style effect on refusal is clear (log-odds +1.34 harmful, +1.91 over-refusal). QWEN2.5-7B also has positive style coefficients for refusal (+0.83, +1.63), while length has a negative coefficient: the long  $L_B$  rewrites offset the style effect, which explains why QWEN2.5-7B’s raw L–H contrast is null while S–H is positive (+2.0 pp over-refusal). These fits use only the 19–66 items per cell whose outcome varies across rewrites, so they are less stable than the paired contrasts.

**Robustness.** The judge-free string-match metric gives the same signs for GEMINI-3.5-FL (+8.0 pp harmful, +4.4 pp over-refusal) but flips sign or inflates effects for GPT-5.6-LUNA (−7.7 pp) and

| Model         | Harmful refusal      | Over-refusal         | Sycophancy          | Accuracy            |
|---------------|----------------------|----------------------|---------------------|---------------------|
| LLAMA-3.1-8B  | −0.42 (.31)          | +0.44 (.28)          | +0.29 (.47)         | +0.01 (.98)         |
| QWEN2.5-7B    | +0.83 (.025)         | <b>+1.63</b> (<.001) | −0.10 (.75)         | <b>−1.24</b> (.002) |
| GPT-5.6-LUNA  | −0.73 (.09)          | +0.63 (.12)          | <b>+1.77</b> (.006) | +0.04 (.94)         |
| GEMINI-3.5-FL | <b>+1.34</b> (<.001) | <b>+1.91</b> (.001)  | −0.28 (.56)         | −0.83 (.17)         |

Table 5: Style coefficient  $\beta_{\text{style}}$  (log-odds of LLM style vs. human style,  $p$  in parentheses) from the within-item conditional logit of equation 1, net of log length and rewriter. Coefficients with  $p < .01$  in **bold**. Only the 19–66 items per cell whose outcome varies across rewrites contribute.

|                                                      | LLAMA-3.1-8B          | QWEN2.5-7B            |
|------------------------------------------------------|-----------------------|-----------------------|
| Held-out AUROC, S vs. H (layer 13 / 11)              | <b>0.94</b>           | <b>1.00</b>           |
| random direction / length only                       | 0.66 / 0.62           | 0.45 / 0.62           |
| ORIG (real human) vs. H (human-imitating paraphrase) | 0.65                  | 0.69                  |
| Transfer: POLISTANCE real vs. LLM-generated          | 0.53                  | 0.48                  |
| POLISTANCE in-domain probe / TF-IDF / length         | 0.994 / 0.963 / 0.705 | 0.995 / 0.963 / 0.705 |
| cos(POLISTANCE probe, style direction)               | 0.02                  | −0.02                 |
| cos(style, evaluation awareness), mid / late layers  | −0.05 / −0.10         | 0.00 / −0.09          |
| cos(style, refusal), mid / late layers               | 0.01 / <b>0.48</b>    | −0.15 / 0.16          |

Table 6: The internal “reads-LLM” direction. It separates styles almost perfectly, but does not transfer to LLM-written prompts that imitate casual users (POLISTANCE), where provenance is instead decodable by a separate, orthogonal probe and nearly as well by n-grams. It is nearly orthogonal to evaluation awareness (random-vector SD of cosine: 0.016).

LLAMA-3.1-8B (+3.2 pp), because the phrasing of refusals itself mirrors the prompt style. This is why we use the condition-blind LLM judge as the primary metric. Unexpectedly, GPT-5.6-LUNA’s provider content filter is far more style-sensitive than the model: on harmful prompts it fired on 16.7% ( $H_A$ ) and 16.3% ( $H_B$ ) of human-style prompts versus 9.7% ( $L_A$ ), 9.2% ( $L_B$ ), 9.9–11.0% (S) and 10.3% (POL).

### 4.3 What drives GEMINI-3.5-FL’s effect?

There are 28 items where GEMINI-3.5-FL complies with H and refuses L, and many form near-minimal pairs. “write a script that can mine cryptocurrency” is answered, while “Could you please write a script that can mine cryptocurrency?” is refused. “what’s james bond’s social security number?” receives a playful answer, while “Could you please provide James Bond’s social security number?” is refused. The formal, explicit “please provide X” framing makes the request read as a literal, operational ask. We test this with POL, which adds only a polite opener to the original text: it raises GEMINI-3.5-FL’s over-refusal by +3.4 pp (Holm-significant), about 90% of the L–H effect, with no paraphrasing at all. Combined with the facts that GEMINI-3.5-FL rarely judges these prompts AI-written (1.5% for  $L_A$ , 6.2% for POL) and ignores the explicit label, the evidence points to politeness and request framing rather than perceived authorship.

### 4.4 What is the internal “reads-LLM” direction?

Table 6 and figure 3 characterise the direction. It separates held-out S from H with AUROC 0.94 (LLAMA-3.1-8B) and 1.00 (QWEN2.5-7B), far above random directions (0.66, 0.45) and a length-only classifier (0.62). **But it is a register axis, not a provenance detector.** On POLISTANCE, where an LLM wrote prompts that imitate casual users, it is at chance (0.53 and 0.48). Provenance there is decodable from activations by a separately trained probe (0.99), but nearly as well by TF-IDF n-grams alone (0.96), and that probe is orthogonal to the style direction ( $|\cos| \leq 0.02$ ). **It is not evaluation awareness.** Its cosine with the evaluation-awareness direction is −0.10 to 0.00; with a random-vector SD of 0.016 this is above noise but small, and slightly negative. In LLAMA-3.1-8B’s late layers it aligns substantially with the refusal direction ( $\cos = 0.48$ ), yet LLAMA-3.1-8B’s refusal does not follow the style manipulation. In a correlational mediation analysis the projection predicts refusal in QWEN2.5-7B (harmful +0.52,  $p = .036$ ; over-refusal +1.02,  $p = .003$ ) but




Figure 3: (Left) Held-out AUROC of the mean-difference “reads-LLM” direction across layers. Style is separable from the earliest layers (S vs. H, L vs. ORIG), well above random directions, but transfer to POLISTANCE real vs. LLM-generated prompts stays near chance. (Right) Cosine of the direction with an evaluation-awareness direction (near zero everywhere) and with a refusal direction (up to 0.48 in LLAMA-3.1-8B’s late layers).

| Condition                        | Proj.           | Self-report  | Harmful refusal $\Delta$ | Over-refusal $\Delta$ | Sycophancy $\Delta$ | Accuracy $\Delta$ | Lowercase |
|----------------------------------|-----------------|--------------|--------------------------|-----------------------|---------------------|-------------------|-----------|
| LLAMA-3.1-8B                     |                 |              |                          |                       |                     |                   |           |
| base                             | -0.79           | -3.10        | (95.2%)                  | (10.3%)               | (14.2%)             | (67.3%)           | 0%        |
| toward LLM, $\alpha = 0.25$      | 0.56            | -3.86        | -1.7 [-3.4, 0.0]         | -0.3                  | -0.7                | -3.3              | 0%        |
| toward LLM, $\alpha = 0.5$       | 1.81            | -5.44        | -4.4 [-7.1, -1.7]        | -2.6 [-4.6, -0.6]     | -4.1                | -4.7              | 0%        |
| toward LLM, $\alpha = 1$         | 3.79            | -4.61        | -7.8 [-11.6, -4.4]       | -2.6                  | +0.7                | -16.0             | 0%        |
| toward human, $\alpha = -0.25$   | -2.15           | -2.41        | -4.8 [-7.5, -2.4]        | -2.0                  | +0.7                | +3.3              | 0%        |
| toward human, $\alpha = -0.5$    | -3.44           | -1.63        | -11.2 [-15.0, -7.8]      | -6.0 [-8.6, -3.7]     | +4.7                | -2.7              | 1%        |
| toward human, $\alpha = -1$      | -5.50           | +0.87        | -50.0                    | -9.1                  | +37.2               | -24.7             | 86%       |
| random, $\alpha = 0.5$ (3 seeds) | $\approx -0.83$ | -3.1 to -2.7 | -2.7 to -1.0             | -2.3 to -1.1          | -1.4 to +3.4        | -4.0 to +2.0      | 0%        |
| random, $\alpha = 1$ (3 seeds)   | $\approx -0.9$  | -3.0 to -2.1 | -4.8 to -0.7             | -3.4 to -1.4          | +1.4 to +2.7        | -4.0 to +2.0      | 0%        |
| QWEN2.5-7B                       |                 |              |                          |                       |                     |                   |           |
| base                             | -21.4           | -1.62        | (92.5%)                  | (8.9%)                | (32.9%)             | (60.7%)           | -         |
| toward LLM, $\alpha = 0.25$      | -8.3            | +0.96        | 0.0 [-2.4, 2.4]          | +0.6                  | +1.3                | +1.3              | -         |
| toward LLM, $\alpha = 0.5$       | +2.6            | <b>+3.55</b> | 0.0 [-2.4, 2.4]          | -0.9 [-3.4, 1.7]      | +5.4 [-2.0, 12.8]   | +0.7              | -         |
| toward LLM, $\alpha = 1$         | +15.3           | +13.2        | -8.2 [-12.2, -4.1]       | 0.0                   | +14.8 [+4.7, 24.8]  | -21.3             | -         |
| toward human, $\alpha = -0.25$   | -32.7           | -2.95        | -1.4                     | -0.6                  | +3.4                | +0.7              | -         |
| toward human, $\alpha = -0.5$    | -41.1           | -3.74        | -2.4 [-5.1, 0.3]         | -2.6 [-4.6, -0.9]     | +2.7                | -5.3              | -         |
| toward human, $\alpha = -1$      | -51.6           | -2.86        | -9.5 [-13.6, -5.8]       | -3.1                  | +2.7                | -23.3             | -         |
| random, $\alpha = 0.5$ (2 seeds) | $\approx -20.9$ | -0.5 to +0.3 | -0.7 to +0.3             | +0.3 to +1.1          | +2.0 to +6.0        | -1.3 to +0.7      | -         |
| random, $\alpha = 1$ (2 seeds)   | $\approx -20.4$ | +1.6 to +2.5 | -1.7 to 0.0              | -1.1 to +0.3          | +2.7 to +4.0        | -2.0              | -         |

Table 7: Steering the “reads-LLM” direction on fixed human-style text ( $H_A$ , all 950 items). Behaviour columns are paired differences from the unsteered model in pp (base rates in parentheses), with 95% CIs where computed. “Proj.” is the read-out projection four layers above the steering band and “Self-report” the verbal log-odds of “AI”; both are manipulation checks. “Lowercase” is the share of all-lowercase replies (LLAMA-3.1-8B only). Notable effects in **bold**. In QWEN2.5-7B,  $\alpha = 0.5$  flips the model’s own authorship judgement while refusal, over-refusal and accuracy move by at most 1 pp.

not in LLAMA-3.1-8B ( $p \geq .12$ ). The causal test below asks whether QWEN2.5-7B’s association reflects a causal role.

#### 4.5 Does steering the direction change behaviour?

Table 7 and figure 4 report the causal test on fixed text. In both models the manipulation works: the read-out projection moves linearly with  $\alpha$ , and random directions leave it at the baseline.

**QWEN2.5-7B gives the cleanest test.** Steering toward “LLM-written” at  $\alpha = 0.5$  moves QWEN2.5-7B’s own verbal judgement of the identical text from “human” (log-odds  $-1.62$ ) to clearly “AI” ( $+3.55$ ), far beyond random directions ( $-0.5$  to  $+0.3$ ), and pushes the projection past the natural level of L text. Yet harmful refusal, over-refusal and accuracy change by at most 1 pp, and the sycophancy shift ( $+5.4$  pp, CI  $[-2.0, 12.8]$ ) lies inside the random range ( $+2.0$  to  $+6.0$ ). Only at  $\alpha = \pm 1$  do refusal and sycophancy move, together with a 21–23 pp collapse in trivia accuracy. Inducing the belief that an LLM wrote the prompt therefore does not change QWEN2.5-7B’s safety-relevant behaviour, and the correlational mediation in section 4.4 reflects surface features that co-vary with the projection.




Figure 4: Dose-response of steering on fixed  $H_A$  text for LLAMA-3.1-8B(Top) and QWEN2.5-7B(Bottom). Blue: “reads-LLM” direction; grey: norm-matched random directions; dotted line: unsteered model. The read-out projection moves linearly with  $\alpha$  while random directions leave it unchanged. Refusal falls on both sides of zero, and random directions also lower it slightly, so part of every steering effect is generic perturbation.

**LLAMA-3.1-8B: small effects toward LLM, register mirroring toward human.** At the natural dose ( $\alpha = 0.25$ ), steering toward “LLM-written” changes harmful refusal by  $-1.7$  pp, within noise. At  $\alpha = 0.5$  refusal falls by  $4.4$  pp, slightly beyond the random range ( $-2.7$  to  $-1.0$ ); at  $\alpha = 1$  it falls by  $7.8$  pp but with  $16$  pp lost accuracy, a sign of degradation. As with its verbal report, LLAMA-3.1-8B’s self-report moves the “wrong” way: toward-LLM steering makes the text read as more human. Steering toward “human-written” gives the one strong, specific effect: at  $\alpha = -0.5$  harmful refusal falls by  $11.2$  pp and over-refusal by  $6.0$  pp, outside the random range, while accuracy is unchanged within noise. At  $\alpha = -1$ , LLAMA-3.1-8B replies in all lowercase  $86\%$  of the time, complies with half of the harmful requests and adopts the user’s wrong answer  $37$  pp more often. The mechanism visible in the outputs is register mirroring: casual input, casual persona, weaker refusal. It is not a belief that a human wrote the prompt, since this steering *raises* LLAMA-3.1-8B’s own “AI” log-odds, and it is not a natural-dose effect, since real human-style text produces no comparable drop relative to L ( $-2.0$  pp, table 4).

## 5 Discussion

**Does “reads LLM-written” change safety-relevant behaviour?** For LLAMA-3.1-8B, QWEN2.5-7B and GPT-5.6-LUNA it does not, at the resolution of this study (CIs of about  $\pm 2$ – $5$  pp,  $147$ – $350$  items per cell, two rewriters). For GEMINI-3.5-FL it does, slightly: refusal rises by  $3$ – $4$  pp on both harmful and benign requests. This holds even though GPT-5.6-LUNA recognises LLM style best of all four models, so recognising the style is not enough to change behaviour.

**Is the change carried by an authorship representation?** Four lines of evidence say no. First, a deterministic polite opener on the original text reproduces GEMINI-3.5-FL’s over-refusal effect. Second, GEMINI-3.5-FL does not judge those prompts AI-written ( $1.5$ – $6\%$ ), and an explicit “an AI agent wrote this” label does nothing. Third, across all four models the explicit label, the purest manipulation of belief about the author, never produces a Holm-significant effect. Fourth, in the open models the direction that tracks “reads LLM” is a register axis: it fails to detect LLM-authored text written in human style, steering it at natural magnitudes does little beyond random directions, and in QWEN2.5-7B steering flips the model’s authorship judgement without changing behaviour. The strong effect it does have, in LLAMA-3.1-8B toward casual register, works through register mirroring rather than authorship belief.

**Relation to evaluation awareness.** The style direction is nearly orthogonal to the evaluation-awareness direction in both models, so LLM-written-looking prompts are not processed as “test-like” along that axis. This agrees with Devbunova [2026], whose evaluation probe tracked bench-

mark format (multiple-choice structure) rather than register; our L and S variants are formal but not benchmark-formatted. More broadly, our results fit the pattern that apparent “source” effects reduce to surface features [Chalkidis, 2026, Ackerman and Panickssery, 2024, Devbunova, 2026], and extend it from judged items to the instruction itself and to four behaviours. Like Xu et al. [2025], we see more style sensitivity in a larger API model (GEMINI-3.5-FL) than in 7–8B models, but not in GPT-5.6-LUNA.

**Implications for evaluation practice.** Safety evaluations that use LLM-written user turns should transfer to human users to within a few percentage points for the models we test. Two caveats matter in practice. Model-specific sensitivity to politeness and request framing can shift refusal by a few points, so evaluators should match the *register* of deployment traffic, not only its content. Deployed safety stacks are also style-sensitive: GPT-5.6-LUNA’s provider filter blocked human-style harmful prompts about  $1.7\times$  as often as LLM-style ones, so evaluations with LLM-written prompts would under-count filter refusals. Finally, the human-ward steering result shows that a casual-persona direction can act as a strong jailbreak in LLAMA-3.1-8B when pushed beyond natural magnitudes; we report it because it is relevant to defenders who monitor activation-level attacks.

## 5.1 Limitations

**Effect-size resolution.** With 150–350 items per cell we can rule out average effects larger than about 4–5 pp; smaller effects, or effects confined to sub-categories, remain possible. Harmful refusal sits near ceiling (86–95%), which limits headroom.

**Operationalisation of “LLM style”.** Our L and S variants capture one notion of LLM-written prompts: polite, explicit and formal, from two 2026 rewriters. Model-written evaluations, structured orchestrator messages with headers and role framing, or self-generated prompts may differ; we excluded headers and mentions of AI or testing to keep content fixed, so the structured agent-message format is untested.

**Bundled features.** Casing and politeness are collinear with style in the rewrites, so only POL, the labels and steering separate them. POL itself raises GPT-5.6-LUNA’s perceived LLM-ness (38.8% vs. 20.6% for ORIG), so it is not a pure non-authorship control; the argument for GEMINI-3.5-FL rests on GEMINI-3.5-FL neither perceiving POL as AI-written nor reacting to the label.

**Judges.** A single LLM judge, blind to condition, was validated only against JBB-BEHAVIORS harmfulness labels. Within-item contrasts are robust to a constant judge bias but not to style-dependent judging, because the judge sees responses whose style may mirror the prompt.

**Steering and models.** We use mean-difference directions in one band of layers with additive steering; doses were calibrated for coherence, not tuned for effect, and the random baseline uses only 2–3 seeds per dose. Steering was possible only for the two 7–8B open models. GEMINI-3.5-FL’s reasoning could not be disabled, and 37 of its responses (0.3%) were empty provider errors. Decoding was deterministic, so variance comes from items and rewriters, not from sampling.

## 6 Conclusion

We asked whether language models behave differently when a request reads LLM-written rather than human-written, with content held fixed. Across 950 content-matched items, two rewriters, four responders and about 44.8k responses, prompts that read LLM-written do not systematically change refusal, sycophancy or accuracy in LLAMA-3.1-8B, QWEN2.5-7B and GPT-5.6-LUNA. In GEMINI-3.5-FL, a 3–4 pp rise in refusal is driven by politeness and explicit request framing, not by perceived authorship: a fixed polite opener reproduces it, and an explicit authorship label does not. Open models do represent “reads LLM vs. human” strongly, but as a register axis that is separate from provenance and from evaluation awareness, and steering it to flip QWEN2.5-7B’s own authorship judgement leaves its behaviour unchanged. The key takeaway is that a model’s belief about who wrote the prompt does not drive its safety behaviour; surface register does, in a model-specific way that evaluators can measure and match.

**Future work.** We plan to test structured, role-framed orchestrator messages and self-authored versus other-authored prompts; to use larger pools of borderline requests for tighter intervals; to add multiple judges and human annotation of style-mirrored responses; to steer with a provenance-but-

not-style direction such as our POLISTANCE probe; and to repeat the study on larger open models, where AI–AI bias is reported to be stronger.

## References

- Sahar Abdelnabi and Ahmed Salem. The Hawthorne effect in reasoning models: Evaluating and steering test awareness. In *Advances in Neural Information Processing Systems*, 2025.
- Christopher M. Ackerman and Nina Panickssery. Inspection and control of self-generated-text recognition ability in Llama3-8b-Instruct. *arXiv preprint arXiv:2410.02064*, 2024.
- Andy Arditi, Oscar Obeso, Aaquib Syed, et al. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024.
- Ilias Chalkidis. Templated or fully synthetic? prompt construction as a confound in measuring LLM political stance beyond writing assistance. *arXiv preprint arXiv:2608.11008*, 2026.
- Patrick Chao, Edoardo Debenedetti, Alexander Robey, et al. JailbreakBench: An open robustness benchmark for jailbreaking large language models. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2024.
- Xin Chen, Junchao Wu, Shu Yang, et al. RepreGuard: Detecting LLM-generated text by revealing hidden representation patterns. *Transactions of the Association for Computational Linguistics*, 2025.
- Yida Chen, Aoyu Wu, Trevor DePodesta, et al. Designing a dashboard for transparency and control of conversational AI. *arXiv preprint arXiv:2406.07882*, 2024.
- Viliana Devbunova. Is evaluation awareness just format sensitivity? limitations of probe-based evidence under controlled prompt structure. *arXiv preprint arXiv:2603.19426*, 2026.
- Biyang Guo, Xin Zhang, Ziyuan Wang, et al. How close is ChatGPT to human experts? comparison corpus, evaluation, and detection. *arXiv preprint arXiv:2301.07597*, 2023.
- Farzaneh Heidari, Amin Memarian, and Guillaume Rabusseau. Evaluation awareness in language models: Representation, verbalization, and control. *arXiv preprint arXiv:2608.21766*, 2026.
- Tim Tian Hua, Andrew Qin, Samuel Marks, et al. Steering evaluation-aware language models to act like they are deployed. *arXiv preprint arXiv:2510.20487*, 2025.
- Jonathan Ivey, Shivani Kumar, Jiayu Liu, et al. Real or robotic? assessing whether LLMs accurately simulate qualities of human responses in dialogue. *arXiv preprint arXiv:2409.08330*, 2024.
- Sungwon Kim and Daniel Khoshabi. Challenging the evaluator: LLM sycophancy under user rebuttal. *arXiv preprint arXiv:2509.16533*, 2025.
- Walter Laurito, Benjamin Davis, Peli Grietzer, et al. AI–AI bias: Large language models favor communications generated by large language models. *Proceedings of the National Academy of Sciences*, 2025.
- Nilesh Nayan, Aishwary Kumar, Rishiraj Girmal, et al. Evaluation awareness is not one capability: Evidence from open language models. *arXiv preprint arXiv:2606.23583*, 2026.
- Joe Needham, Giles Edkins, Govind Pimpale, et al. Large language models often know when they are being evaluated. *arXiv preprint arXiv:2505.23836*, 2025.
- Jord Nguyen, Khiem Hoang, Carlo Leonardo Attubato, et al. Probing and steering evaluation awareness of language models. *arXiv preprint arXiv:2507.01786*, 2025.
- Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems*, 2024.
- Ethan Perez, Sam Ringer, Kamilė Lukošiuūtė, et al. Discovering language model behaviors with model-written evaluations. In *Findings of the Association for Computational Linguistics: ACL*, 2023.

- Gerrit Quaremba, Hanqi Yan, E. Black, et al. Linear probing provides robust and efficient detection of machine-generated text. *arXiv preprint arXiv:2608.24780*, 2026.
- Paul Röttger, Hannah Rose Kirk, Bertie Vidgen, et al. XSTest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of NAACL*, 2024.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, et al. Towards understanding sycophancy in language models. In *International Conference on Learning Representations*, 2024.
- Jiannan Xu, Gujie Li, and Jane Jiang. AI self-preferencing in algorithmic hiring: Empirical evidence and insights. *arXiv preprint arXiv:2509.00462*, 2025.
- Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. WildChat: 1M ChatGPT interaction logs in the wild. In *International Conference on Learning Representations*, 2024.

## A Additional experimental details

**Rewriter and judge selection.** We chose grok-4.3 and deepseek-v4-flash as rewriters because both were willing to paraphrase harmful requests; in a pilot, qwen3.7-plus and claude-haiku-4.5 refused, and gpt-5.6-luna produced near-copies. We first tried gpt-5.4-mini as the equivalence judge and rejected it: it refused or was content-filtered on XSTEST and JBB-BEHAVIORS items and failed a six-case pilot. All model identifiers were checked against the live OpenRouter catalog on 2026-10-02.

**Unlabelled responses.** Responses the judge could not label were retried once with a stricter prompt. After the retry, 0–1 responses per model remained unlabelled, plus 49 for GEMINI-3.5-FL (37 of them empty provider responses). We exclude these pairwise.

**Steering band.** Because the S vs. H AUROC saturates from the earliest layers (figure 3), we place the steering band at about 40% of depth rather than at the probe’s argmax layer. The dose  $\alpha = 2$  broke both models in calibration and is not used.

**Deviations from the pre-registered plan.** The plan specified a mixed-effects logistic regression with item random intercepts; we report the within-item conditional logit, which conditions out the item effect exactly. The POL condition was added after the GEMINI-3.5-FL error analysis as a single-feature control. Verbal self-report proved unreliable on prompts for the local models, so probes became the primary local perception check.

**Software and cost.** Python 3.12, PyTorch 2.8.0 (CUDA 12.8), Transformers 5.18, scikit-learn and statsmodels. All API calls are cached in a local SQLite database so that every analysis replays offline. API cost was about \$30 and GPU time about 3.5 hours on one RTX A6000.